



UBLK

HIGH PERFORMANCE GENERIC USERSPACE BLOCK DEVICE

Ming Lei <ming.lei@redhat.com>, Red Hat

What is UBLK

- High performance generic userspace block device
- Goals:
 - High performance
 - Expose generic block device, and support all kinds of block/queue settings/parameters
 - Move all block IO logic in userspace
 - Implement userspace target/backend easily



Motivations

- can be written many programming languages.
- can use libraries that are not available in the kernel.
- can be debugged with tools familiar to application developers.
- Crashes do not kernel panic the machine.
- Bugs are likely to have a lower security impact than bugs in kernel
- can be installed and updated independently of the kernel.
- can be used to simulate block device easily with user specified parameters/setting for test/debug purpose

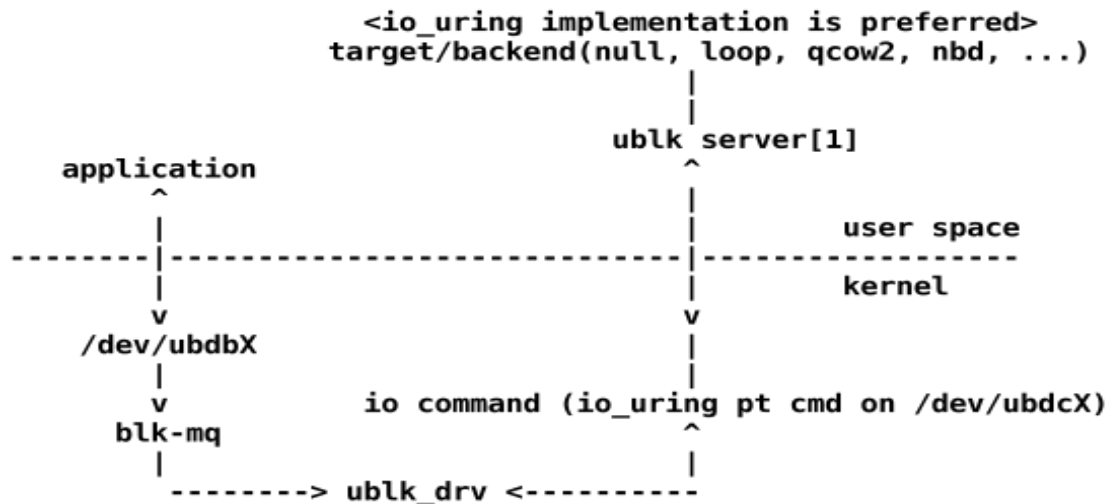


Background

- NBD, merged to linux kernel 2.1.15 in 1997
 - expose nbd device node, socket communication
- VDUSE, merged to linux kernel 5.5 in 2021
 - expose as virtio_blk, io command via traditional read/write on char device
- UBLK, merged to linux kernel 6.0 in 2022, io command via io_uring pt cmd
- BDUS: 2021 <https://dl.acm.org/doi/10.1145/3456727.3463768>
- BUSE: 2021 <https://github.com/acozzette/BUSE>
- DM-USER: 2020 <https://lwn.net/Articles/838986/>
- More...



UBLK framework



[1] ublk server is term for generic userspace implementation, and ublksrv is one such implementation

[2] ublksrv: <https://github.com/ming1/ubdsrv>



UBLK framework

- ublk drv
 - in linux kernel v6.0
- ublk server: ublksrv
 - libublksrv
 - ublksrv generic target/backend
 - ublksrv target/backend

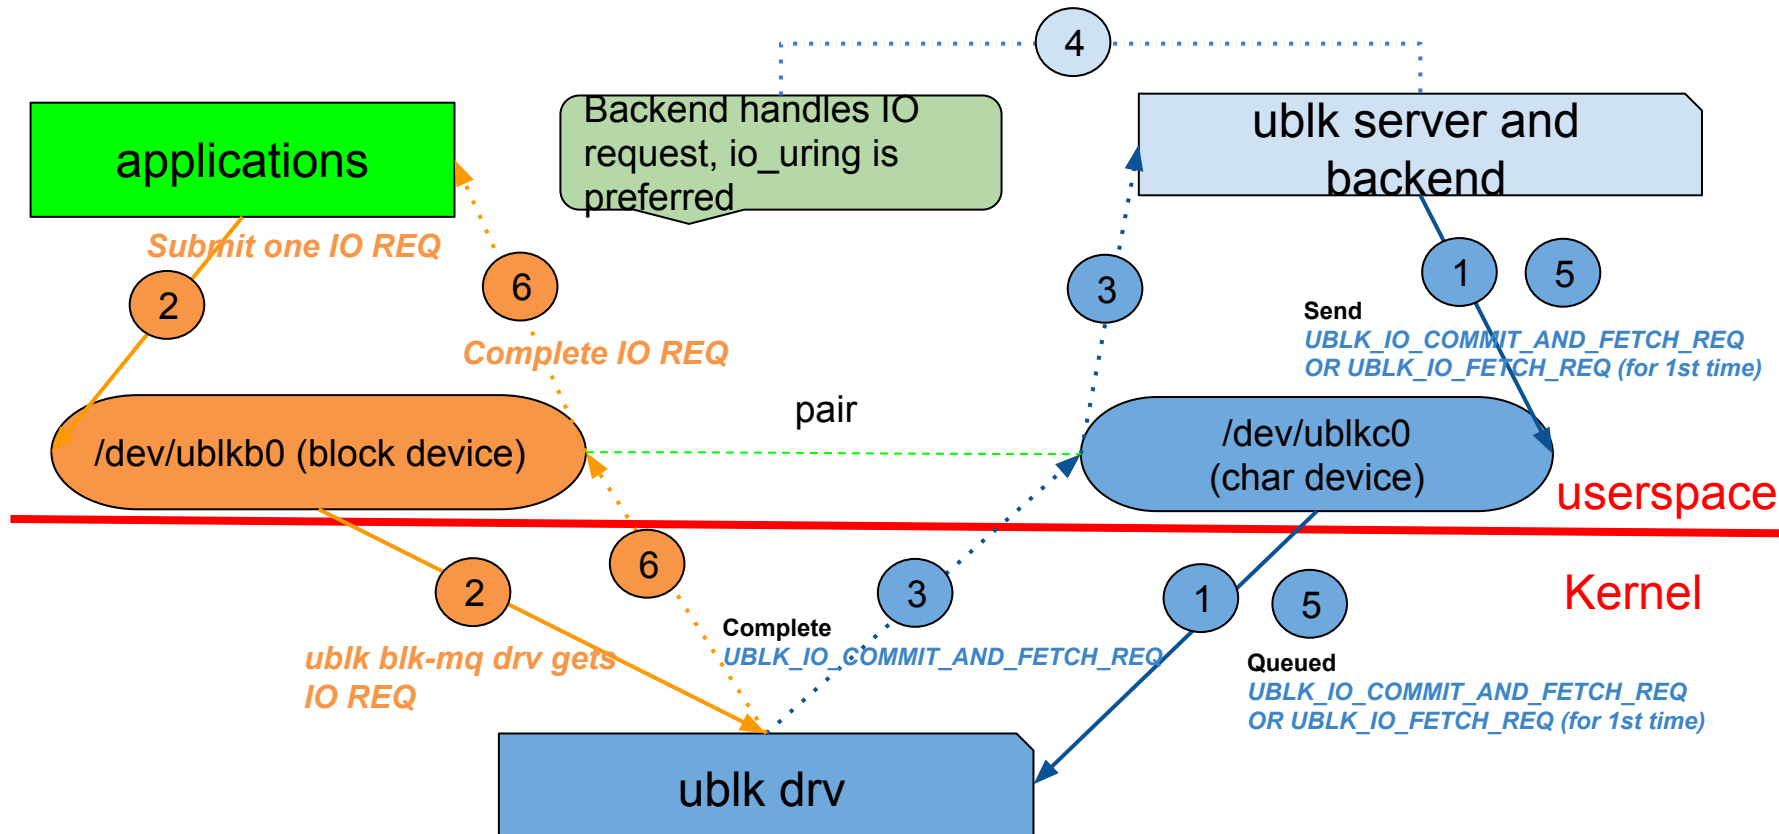


IO command communication

- IO descriptor
 - each IO has unique per-queue tag
 - IO descriptor is written to shared/mmapped area which can be indexed by io tag, read-only for ublk server, and write-only for ublk drv
- UBLK_IO_FETCH_REQ(io_uring pt cmd)
 - sent once from ublk server for setting up IO communication
- UBLK_IO_COMMIT_AND_FETCH_REQ (io_uring pt cmd)
 - When ublk IO req comes, the issued *_FETCH_REQ is completed
 - After the IO is handled by ublk server, this command is issued to ublk drv for both committing previous IO result and start to fetch new request



IO command communication



ublkdrv

- Userspace part:
 - <https://github.com/ming1/ubdrv>
 - in Fedora Rawhide
 - add / del / list / recovery device
 - libublkdrv: provide generic interface for standalone target development, such nbdublk,
 - IO target/backend implementation
- Supported targets(early stage)
null, loop, qcow2(basic read/write)
- Preferred target io handling
io_uring for getting top performance



UBLK performance

- ublk-loop: IOPS is close to kernel loop with `-directio=on`
 - <https://lwn.net/Articles/903855/>
 - <https://lore.kernel.org/all/20220713140711.97356-1-ming.lei@redhat.com/>
- ublk vs. qemu-nbd, by comparing qcow2 target
 - > 2~3X IOPS in random IO test
 - <https://lore.kernel.org/lkml/Yza1u1KfKa7ycQm0@T590/>
- ublk vs. vduse:
 - 1job 1 io depth: ½ latency of vduse over null_blk
 - 4job 128 io depth: ~3X IOPS of vduse
 - <https://lore.kernel.org/lkml/50827796-af93-4af5-4121-dc13c31a67fc@linux.alibaba.com/>



Why does UBLK perform so well

- High performance io uring passthrough command
 - io_uring pt cmd is proved as efficient, even more than io_uring over block IO
 - IO command is submitted beforehand, minimize io command forward latency
 - IO command multiplexing: one command covers both result committing and fetching new req
- target/backend IO handling by io_uring too
 - share same io_uring context, maximize io batching in single syscall
- IO handle efficiently
 - each IO has its unique tag, submit io command/allocate resource beforehand
 - work together with per-IO stackless coroutine, minimize context switch and maximize IO parallelization
 - meantime simplify IO handling development




Future development

- Container-ware ublk
 - to be unprivileged, actually both io command submission & completion & handling are done in user task
- Zero copy for big chunk IO
 - is it possible to avoid the single pages copy for big chunk IO?
- All kinds of performance improvement
 - sequential big chunk IO has improvement space, get user pages latency
- Cross platform
 - io_uring is supported by windows 11
- More targets/backends
 - nbd, zoned, compressed, rbd, iscsi, nvme-tcp, ...
 - make full use of io_uring's high performance advantage



Questions

- Email / github
 - ming.lei@redhat.com
 - <https://github.com/ming1/ubdsrv.git>
- Welcome to participate in ublk development





THANK YOU



plus.google.com/+RedHat



facebook.com/redhatinc



linkedin.com/company/red-hat



twitter.com/RedHatNews



youtube.com/user/RedHatVideos